

Exploring the weather dataset

ANALYZING POLICE ACTIVITY WITH PANDAS



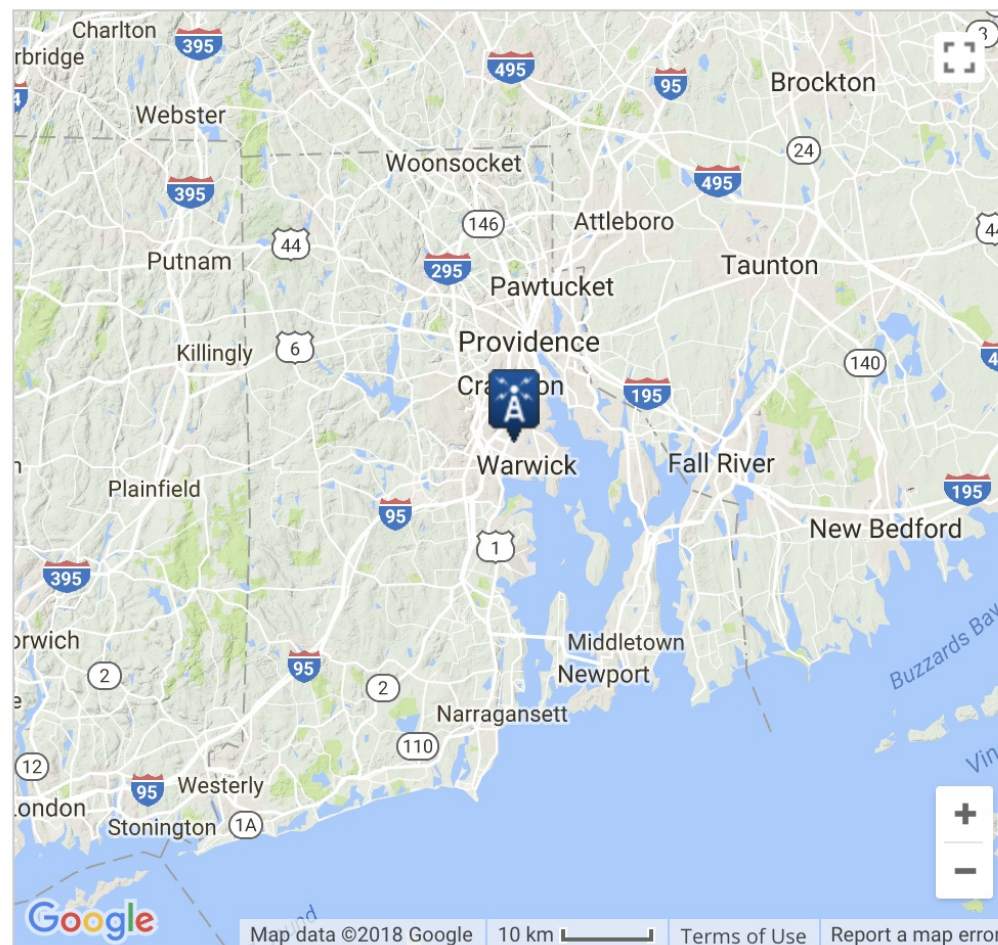
Kevin Markham
Founder, Data School

Introduction to the dataset



STATION DETAILS	
Name	PROVIDENCE, RI US
Network:ID	GHCND:USW00014765
Latitude/Longitude	41.7225°, -71.4325°
Elevation	16.8 m

PERIOD OF RECORD	
Start Date ¹	1942-08-01
End Date ¹	2018-04-21
Data Coverage ²	93%



```
weather = pd.read_csv('weather.csv')
weather.head(3)
```

	STATION	DATE	TAVG	TMIN	TMAX	AWND	WSF2	WT01	WT02	
0	USW00014765	2005-01-01	44.0	35	53	8.95	25.1	1.0	NaN	
1	USW00014765	2005-01-02	36.0	28	44	9.40	14.1	NaN	NaN	
2	USW00014765	2005-01-03	49.0	44	53	6.93	17.0	1.0	NaN	
	...									
	WT11	WT13	WT14	WT15	WT16	WT17	WT18	WT19	WT21	WT22
0	...	NaN	1.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	...	NaN	NaN	NaN	NaN	1.0	NaN	1.0	NaN	NaN
2	...	NaN	1.0	NaN	NaN	1.0	NaN	NaN	NaN	NaN

- TAVG , TMIN , TMAX : Temperature
- AWND , WSF2 : Wind speed
- WT01 ... WT22 : Bad weather conditions

Examining the wind speed

```
weather[['AWND', 'WSF2']].head()
```

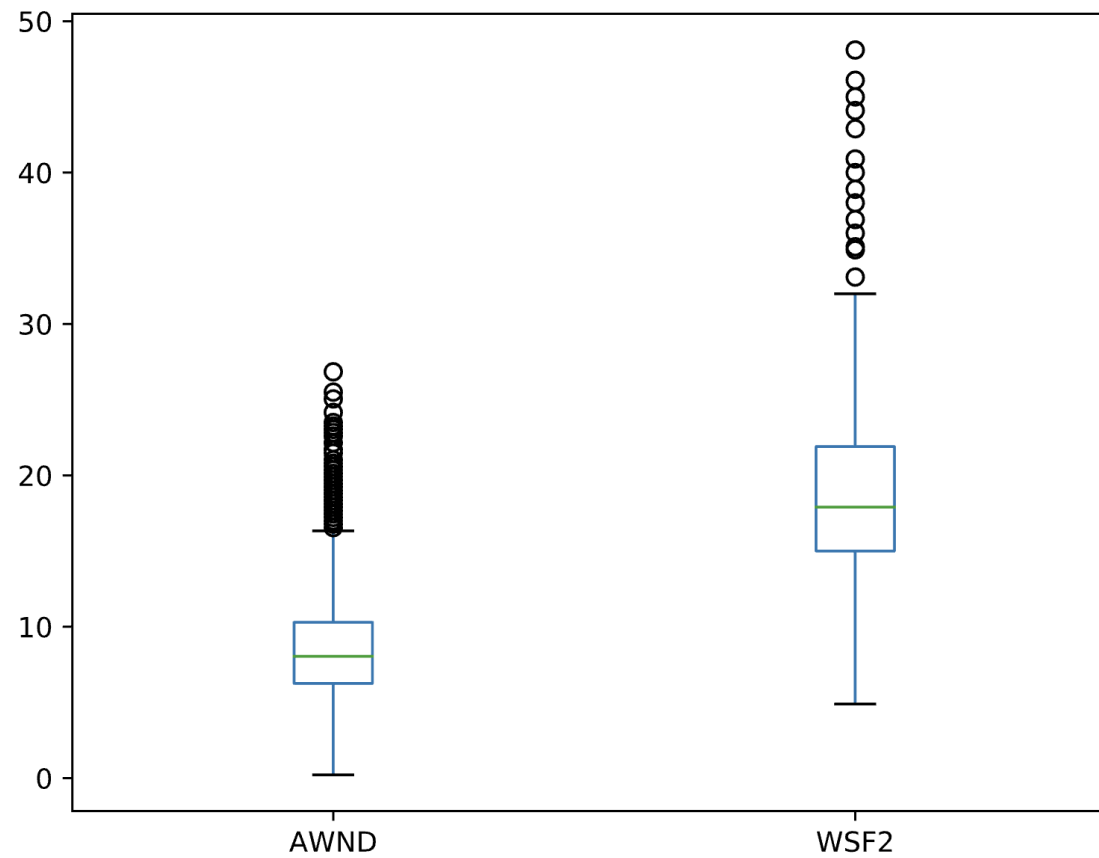
	AWND	WSF2
0	8.95	25.1
1	9.40	14.1
2	6.93	17.0
3	6.93	16.1
4	7.83	17.0

```
weather[['AWND', 'WSF2']].describe()
```

	AWND	WSF2
count	4017.000000	4017.000000
mean	8.593707	19.274782
std	3.364601	5.623866
min	0.220000	4.900000
25%	6.260000	15.000000
50%	8.050000	17.900000
75%	10.290000	21.900000
max	26.840000	48.100000

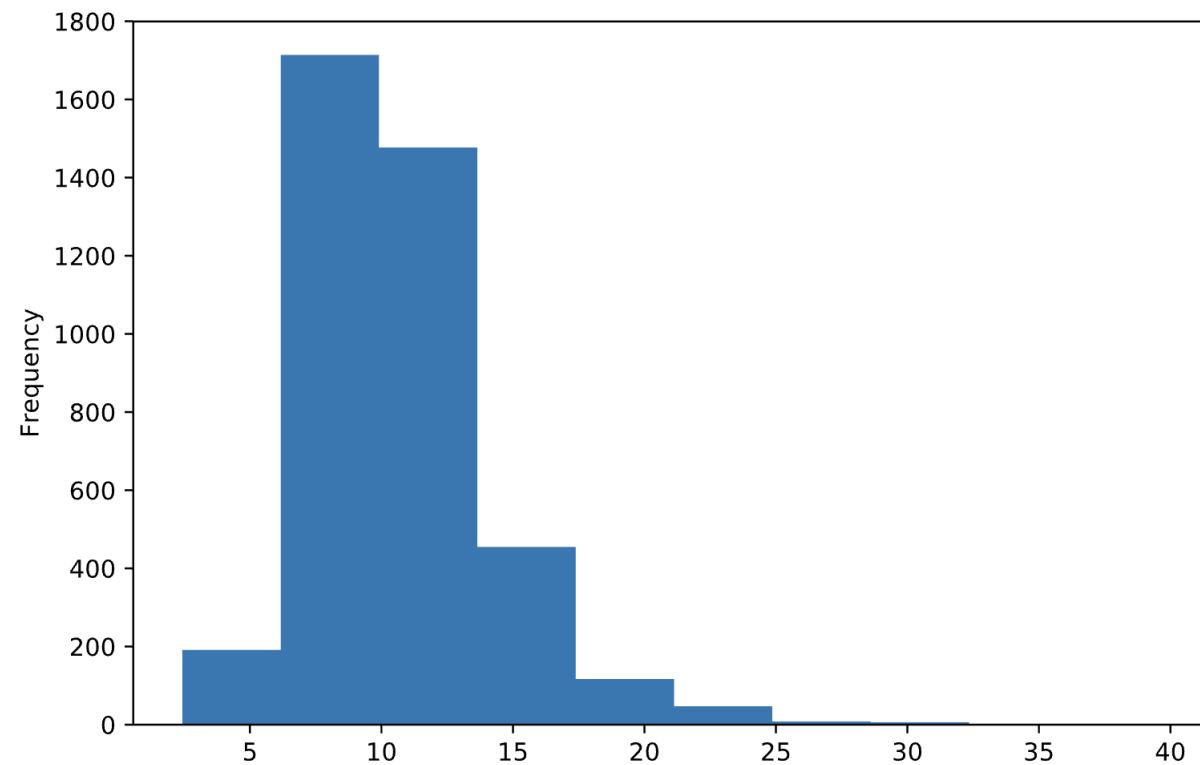
Creating a box plot

```
weather[['AWND', 'WSF2']].plot(kind='box')  
plt.show()
```



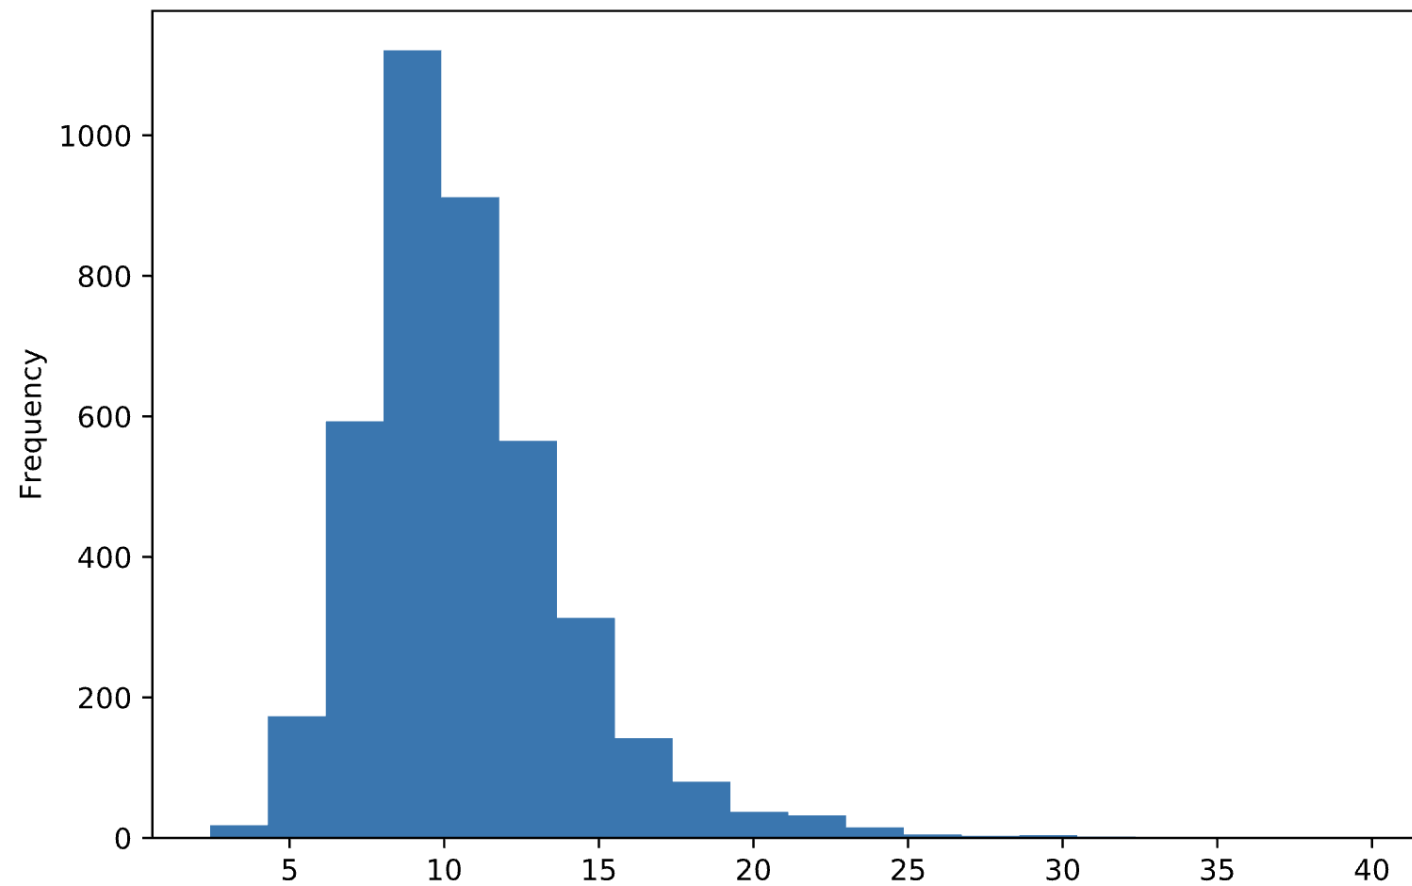
Creating a histogram (1)

```
weather['WDIFF'] = weather.WSF2 - weather.AWND  
weather.WDIFF.plot(kind='hist')  
plt.show()
```



Creating a histogram (2)

```
weather.WDIFFF.plot(kind='hist', bins=20)  
plt.show()
```



Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Categorizing the weather

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham
Founder, Data School

Selecting a DataFrame slice (1)

```
weather.shape
```

```
(4017, 28)
```

```
weather.columns
```

```
Index(['STATION', 'DATE', 'TAVG', 'TMIN', 'TMAX', 'AWND',  
      'WSF2', 'WT01', 'WT02', 'WT03', 'WT04', 'WT05',  
      'WT06', 'WT07', 'WT08', 'WT09', 'WT10', 'WT11',  
      'WT13', 'WT14', 'WT15', 'WT16', 'WT17', 'WT18',  
      'WT19', 'WT21', 'WT22', 'TDIFF'], dtype='object')
```

Selecting a DataFrame slice (2)

```
temp = weather.loc[:, 'TAVG' : 'TMAX' ]
```

```
temp.shape
```

```
(4017, 3)
```

```
temp.columns
```

```
Index(['TAVG', 'TMIN', 'TMAX'], dtype='object')
```

DataFrame operations

```
temp.head()
```

	TAVG	TMIN	TMAX
0	44.0	35	53
1	36.0	28	44
2	49.0	44	53
3	42.0	39	45
4	36.0	28	43

```
temp.sum(axis='columns').head()
```

0	132.0
1	108.0
2	146.0
3	126.0
4	107.0

```
temp.sum()
```

TAVG	63884.0
TMIN	174677.0
TMAX	246116.0

Mapping one set of values to another

```
ri.stop_duration.unique()
```

```
array(['0-15 Min', '16-30 Min', '30+ Min'], dtype=object)
```

```
mapping = {'0-15 Min': 'short',  
          '16-30 Min': 'medium',  
          '30+ Min': 'long'}  
  
ri['stop_length'] = ri.stop_duration.map(mapping)  
  
ri.stop_length.dtype
```

```
dtype('O')
```

Changing data type from object to category (1)

```
ri.stop_length.unique()
```

```
array(['short', 'medium', 'long'], dtype=object)
```

- Category type stores the data more efficiently
- Allows you to specify a logical order for the categories

```
ri.stop_length.memory_usage(deep=True)
```

```
8689481
```

Changing data type from object to category (2)

```
cats = ['short', 'medium', 'long']
```

```
ri['stop_length'] = ri.stop_length.astype('category',  
                                          ordered=True,  
                                          categories=cats)
```

```
ri.stop_length.memory_usage(deep=True)
```

```
3400602
```

Using ordered categories (1)

```
ri.stop_length.head()
```

```
stop_datetime  
2005-01-04 12:55:00      short  
2005-01-23 23:15:00      short  
2005-02-17 04:15:00      short  
2005-02-20 17:15:00    medium  
2005-02-24 01:20:00      short  
Name: stop_length, dtype: category  
Categories (3, object): [short < medium < long]
```


Using ordered categories (2)

```
ri[ri.stop_length > 'short'].shape
```

```
(16959, 16)
```

```
ri.groupby('stop_length').is_arrested.mean()
```

```
stop_length
short      0.013654
medium     0.093595
long       0.261572
Name: is_arrested, dtype: float64
```

Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Merging datasets

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham

Founder, Data School

```
apple
```

```
          date  time  price
date_and_time
2018-02-14 09:30:00 2/14/18 9:30 163.04
2018-02-14 16:00:00 2/14/18 16:00 167.37
2018-02-15 09:30:00 2/15/18 9:30 169.79
2018-02-15 16:00:00 2/15/18 16:00 172.99
```

```
apple.reset_index(inplace=True)
```

```
apple
```

```
   date_and_time  date  time  price
0 2018-02-14 09:30:00 2/14/18 9:30 163.04
1 2018-02-14 16:00:00 2/14/18 16:00 167.37
2 2018-02-15 09:30:00 2/15/18 9:30 169.79
3 2018-02-15 16:00:00 2/15/18 16:00 172.99
```

Preparing the second DataFrame

```
high_low
```

```
   DATE    HIGH    LOW
0  2/14/18  167.54  162.88
1  2/15/18  173.09  169.00
2  2/16/18  174.82  171.77
```

```
high = high_low[['DATE', 'HIGH']]
high
```

```
   DATE    HIGH
0  2/14/18  167.54
1  2/15/18  173.09
2  2/16/18  174.82
```

Merging the DataFrames

```
apple_high = pd.merge(left=apple, right=high,  
                       left_on='date', right_on='DATE',  
                       how='left')
```

- `left=apple` : Left DataFrame
- `right=high` : Right DataFrame
- `left_on='date'` : Key column in left DataFrame
- `right_on='DATE'` : Key column in right DataFrame
- `how='left'` : Type of join

apple_high

	date_and_time		date	time	price	DATE	HIGH
0	2018-02-14	09:30:00	2/14/18	9:30	163.04	2/14/18	167.54
1	2018-02-14	16:00:00	2/14/18	16:00	167.37	2/14/18	167.54
2	2018-02-15	09:30:00	2/15/18	9:30	169.79	2/15/18	173.09
3	2018-02-15	16:00:00	2/15/18	16:00	172.99	2/15/18	173.09

apple

	date_and_time		date	time	price
0	2018-02-14	09:30:00	2/14/18	9:30	163.04
1	2018-02-14	16:00:00	2/14/18	16:00	167.37
2	2018-02-15	09:30:00	2/15/18	9:30	169.79
3	2018-02-15	16:00:00	2/15/18	16:00	172.99

high

	DATE	HIGH
0	2/14/18	167.54
1	2/15/18	173.09
2	2/16/18	174.82

Setting the index

```
apple_high.set_index('date_and_time', inplace=True)  
apple_high
```

	date	time	price	DATE	HIG
date_and_time					
2018-02-14 09:30:00	2/14/18	9:30	163.04	2/14/18	167.5
2018-02-14 16:00:00	2/14/18	16:00	167.37	2/14/18	167.5
2018-02-15 09:30:00	2/15/18	9:30	169.79	2/15/18	173.0
2018-02-15 16:00:00	2/15/18	16:00	172.99	2/15/18	173.0

Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Does weather affect the arrest rate?

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham
Founder, Data School

Driver gender and vehicle searches

```
ri.search_conducted.mean()
```

```
0.0382153092354627
```

```
ri.groupby('driver_gender').search_conducted.mean()
```

```
driver_gender
F    0.019181
M    0.045426
```

```
ri.groupby(['violation', 'driver_gender']).search_conducted.mean()
```

violation	driver_gender	
Equipment	F	0.039984
	M	0.071496
Moving violation	F	0.039257
	M	0.061524
Other	F	0.041018
	M	0.046191
Registration/plates	F	0.054924
	M	0.108802
Seat belt	F	0.017301
	M	0.035119
Speeding	F	0.008309
	M	0.027885

```
search_rate = ri.groupby(['violation',  
                          'driver_gender']).search_conducted.mean()  
  
search_rate
```

```
violation      driver_gender  
Equipment      F            0.039984  
               M            0.071496  
Moving violation F            0.039257  
               M            0.061524  
...           ...           ...
```

```
type(search_rate)  
type(search_rate.index)
```

```
pandas.core.series.Series  
pandas.core.indexes.multi.MultiIndex
```

violation	driver_gender	
Equipment	F	0.039984
	M	0.071496
Moving violation	F	0.039257
	M	0.061524
...

```
search_rate.loc['Equipment']
```

```
driver_gender
F    0.039984
M    0.071496
```

```
search_rate.loc['Equipment', 'M']
```

```
0.07149643705463182
```

Converting a multi-indexed Series to a DataFrame

```
search_rate.unstack()
```

```
driver_gender      F      M
violation
Equipment      0.039984  0.071496
Moving violation  0.039257  0.061524
Other           0.041018  0.046191
...              ...      ...
```

```
type(search_rate.unstack())
```

```
pandas.core.frame.DataFrame
```

Converting a multi-indexed Series to a DataFrame

```
ri.pivot_table(index='violation',  
                columns='driver_gender',  
                values='search_conducted')
```

```
driver_gender      F      M  
violation  
Equipment      0.039984  0.071496  
Moving violation  0.039257  0.061524  
Other          0.041018  0.046191  
...
```


Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Conclusion

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham
Founder, Data School

Stanford Open Policing Project



Thank you!

ANALYZING POLICE ACTIVITY WITH PANDAS