



✓ **Congratulations! You passed!**
TO PASS 80% or higher

Keep Learning

GRADE
87.50%

You have new skill scores!

Great job! Keep learning and making progress in your courses to increase your skill scores.[See skills](#)

Machine Learning 273

Computer Programming 232

Week 3 Quiz

LATEST SUBMISSION GRADE
87.5%

1. Why does sequence make a large difference when determining semantics of language?

1 / 1 point

- ☐ Because the order of words doesn't matter
- ☐ It doesn't
- ☐ Because the order in which words appear dictate their meaning
- ☒ Because the order in which words appear dictate their impact on the meaning of the sentence

✓ Correct

2. How do Recurrent Neural Networks help you understand the impact of sequence on meaning?

1 / 1 point

- ☐ They don't
- ☐ They look at the whole sentence at a time
- ☒ They carry meaning from one cell to the next
- ☐ They shuffle the words evenly

✓ Correct

3. How does an LSTM help understand meaning when words that qualify each other aren't necessarily beside each other in a sentence?

1 / 1 point

- ☒ Values from earlier words can be carried to later ones via a cell state
- ☐ They shuffle the words randomly
- ☐ They don't
- ☐ They load all words into a cell state

✓ Correct

4. What keras layer type allows LSTMs to look forward and backward in a sentence?

1 / 1 point

- ☐ Unilateral
- ☐ Bothdirection
- ☒ Bidirectional
- ☐ Bilateral

✓ Correct

5. What's the output shape of a bidirectional LSTM layer with 64 units?

1 / 1 point

- ☐ (None, 64)
- ☒ (None, 128)
- ☐ (128, None)
- ☐ (128, 1)

✓ Correct

6. When stacking LSTMs, how do you instruct an LSTM to feed the next one in the sequence?

1 / 1 point

- ☐ Do nothing, TensorFlow handles this automatically
- ☐ Ensure that they have the same number of units
- ☐ Ensure that return_sequences is set to True on all units
- ☒ Ensure that return_sequences is set to True only on units that feed to another LSTM

✓ Correct

7. If a sentence has 120 tokens in it, and a Conv1D with 128 filters with a Kernal size of 5 is passed over it, what's the output shape?

1 / 1 point

- ☒ (None, 116, 128)
- ☐ (None, 120, 124)
- ☐ (None, 120, 128)
- ☐ (None, 116, 124)

✓ Correct

8. What's the best way to avoid overfitting in NLP datasets?

0 / 1 point

- ☐ Use LSTMs
- ☒ Use GRUs
- ☐ Use Conv1D
- ☐ None of the above

! Incorrect