Congratulations! You passed!

Grade received 100% To pass 80% or higher

Go to next item

Introduction to Model Serving

Latest Submission Grade 100%

1. What are the three key components we should consider when serving an ML Model in a production environment? (Select all that apply) — 1 / 1 point

- [x] An interpreter

  Correct
  Right on track! An Interpreter encapsulates a pre-trained model in which operations are executed for inference.

- [x] Input Data

  Correct
  You've got it! The model executed on-device makes predictions based on the input data.

- [ ] An orchestrator

- [x] A model

  Correct
  Correct! Providing the algorithm and training the ML model is the first step towards putting it into production.

2. What happens after a while in operation to an offline-trained model dealing with new real-live data? — 1 / 1 point

- ( ) The model abruptly forgets all previously learned information.
- (•) The model becomes stale.
- ( ) The model adapts to new patterns.

  Correct
  Good job! The model performance deteriorates to the point of the model not being any longer fit for purpose. This phenomenon is called model decay and should be carefully monitored.

3. In applications that are not user-facing, is throughput more critical than latency for customer satisfaction? — 1 / 1 point

- ( ) No, because users might complain that the app is too slow.
- (•) Yes, in this case, we are concerned with maximizing throughput with the lowest CPU usage.

  Correct
  Correct! Latency is not a key concern for back-end services.

4. Nowadays, developers aim to minimize latency and maximize throughput in customer-facing applications. However, in doing so, infrastructure scales and costs increase. So, what strategies can developers implement to balance cost and customer satisfaction? (Select all that apply) — 1 / 1 point

- [x] GPU sharing

  Correct
  Nailed it! This strategy reduces the cost of GPU-accelerated computing.

- [x] Multi-model serving

  Correct
  Yes! This approach scales back infrastructure.

- [x] Optimizing inference models

  Correct
  Right on track! Optimization modifies a model to handle a higher load, reducing costs as a result.

- [ ] Stress testing

- [ ] Stress testing