! **Try again once you are ready**
TO PASS 80% or higher

Try again

GRADE
40%

Visualization and Interpretation

LATEST SUBMISSION GRADE
40%

1. Consider the following code for Class Activation Maps. Which layer(s) of the *model* do we choose as outputs to draw out the class activation map ? Check all that apply. **0.75 / 1 point**

- [x] The layer which performs *classification* on the model

✓ **Correct**
Correct!

- [ ] The layer which feeds *input* to the model
- [ ] The layer which holds the extracted *features* in the model
- [ ] The layer which performs *concatenation* in the model

You didn't select all the correct answers

2. To compute the Class Activation Map you _____. **1 / 1 point**

- ( ) Take the dot product of the weights associated with the prediction and the output of the classification vector.
- ( ) Take the dot product of the features and the output of the classification vector.
- (•) Take the dot product of the features associated with the prediction on the image, with the weights that come from the last global average pooling layer.

✓ **Correct**
Correct!

3. In a Salience map you get to see parts of the image the model was paying attention to when deciding what class to assign to the image. **0 / 1 point**

- (•) True
- ( ) False

! **Incorrect**
Salience Map is a representation of every pixel in the image in a way that makes sense for your particular image. E.g, what makes a cat a *cat*(the way the ears are, the eyes, the paws)

4. In Saliency Maps, the pixels that most impact the final classification are found by looking at the gradients of the final layers to see which ones had the steepest curve, and figure out their location and plot them on the original image. **0 / 1 point**

- ( ) True
- (•) False

! **Incorrect**
Incorrect! The above statement is true.

5. Which of the following statements are *not true* about GradCAM? Check all that apply. **0.25 / 1 point**

- [ ] The negative values in the *heatmap*of the gradCAM are kept as they enhance the performance and accuracy of the gradCAM.
- [x] The gradients of the loss are computed with respect to the selected layer's output and averaged out across all feature maps.

! **This should not be selected**
Incorrect! This statement is true.

- [ ] You stack the filter outputs on the final layer into a heatmap by calculating the mean of those values.

- [ ] The *model* built to perform the task uses the *last two* layers of the *original* model as the *outputs*.